

Workshop 2 – RESCUER MSCA DOCTORAL NETWORK 2024-2028
Universidad de Zaragoza

GPU programming (II)

Task 2: CUDA kernel implementation

Mario Morales Hernández

mmorales@unizar.es

Module 4 – GPU programming (II)

Task 2 - CUDA kernel implementation

The second day focuses on implementing CUDA kernels for key functions in the shallow water equations solver. The goal is to convert CPU-based functions into efficient GPU kernels while ensuring correctness and performance. The first part will focus on flux computation functions (*h_compute_x_fluxes* and *h_compute_y_fluxes*), while the second part will address wet/dry cell checking (*h_wet_dry_x* and *h_wet_dry_y*) and managing host-device memory transfers for functions that remain on the CPU.

Key tasks:

1. Convert *h_compute_x_fluxes* and *h_compute_y_fluxes*:

- Use the flattened loops from Task 1 to implement CUDA kernels.
- Add the `__global__` keyword to define the kernels. Use the same macro used in Task 1.
- Implement thread indexing using `blockIdx.x`, `threadIdx.x`, and `blockDim.x`.
- Implement the kernel call, paying attention to the grid dimension and `blockSize`.

2. Convert *h_wet_dry_x* and *h_wet_dry_y*:

- Use the flattened loops from Task 1 to implement CUDA kernels.
- Add the `__global__` keyword to define the kernels. Use the same macro used in Task 1.
- Implement thread indexing using `blockIdx.x`, `threadIdx.x`, and `blockDim.x`.
- Implement the kernel call, paying attention to the grid dimension and `blockSize`.

3. Keep certain functions on the CPU: Leave *h_compute_flow_time_step_2D* and *h_check_depth_positivity* on the CPU.

- Use `cudaMemcpy` to transfer data between the host and device as needed.
- Implement proper memory transfers:
- Identify variables required for CPU functions and transfer them efficiently between the host and device.

4. Modify *swe2d.c* to update the kernel calls for the converted subroutines (*h_compute_x_fluxes*, *h_compute_y_fluxes*, *h_wet_dry_x*, and *h_wet_dry_y*). Use the same macro pattern from Day 1 to keep the CPU and GPU versions clean and maintainable. Pay attention to the kernel call configuration, ensuring the correct grid and block dimensions. Use the `nThreads` macro for the block size.